

Problem 1

MDP Formulation

I've decided to formulate the problem as a finite horizon undiscounted MDP problem. There are 5 rounds, and 4 decisions to be made (the last digit is forced). The state consists of a mask (ie boolean vector) indicating which positions are still available (we will define r to be the number of rounds of the game, and denote this mask as $\mathbf{m} \in \mathbb{R}^{r \times 1}$) and the currently observed digit d . The round number t is implicit in the state as $t = \sum_p m_p$.

Note that I have **not** included the partially filled number in the state. It's easier to bake this into the reward structure instead of carrying it around in the state.

On that note, I'll define the reward structure, and some of my reasons for the setup. The reward function is simply the immediate value of placing digit d in position p , which is $d \cdot 10^p$. This feels natural, intuitive, and doesn't require additional bookkeeping nor recalculation at the completion of the game.

Alternatively, to maintain the order of the vector with the lexicographic and familiar ordering, I could define the reward as $d \cdot 10^{r-1-p}$

States. At the point of each decision in round $t \in [r]$, the state is

$$s_t = (\mathbf{m}, d)$$

Like I touched upon in the intro, the state consists of a mask $\mathbf{m} \in \{0, 1\}^{r \times 1}$ indicating which positions are filled ($m_p = 1$ if position p is filled, else 0), and the currently observed digit $d \in \{0, \dots, 9\}$.

Actions. From state (\mathbf{m}, d) , choose any empty position:

$$\mathcal{A}(\mathbf{m}) = \{p \in \{0, \dots, r-1\} : m_p = 0\}.$$

Transition model. Let's define the helper function $\text{fill}(\mathbf{m}, p)$ that returns the mask with position p set to 1:

$$\text{fill}(\mathbf{m}, p) = (m_0, \dots, m_{p-1}, 1, m_{p+1}, \dots, m_{r-1}).$$

After taking action $a = p$ (placing digit d at position p), the next state is

$$\mathbf{m}' = \text{fill}(\mathbf{m}, p), \quad d' \sim \text{Unif}\{0, \dots, 9\}$$

For each $d' \in \{0, \dots, 9\}$,

$$\Pr\left(\text{fill}(\mathbf{m}, p), d'\right) \mid (\mathbf{m}, d), a = p = \frac{1}{10},$$

and all other next states have probability 0.

In plain English, the next state is the same as the current state except that position p is now filled, and the next digit is drawn uniformly from $\{0, \dots, 9\}$.

Rewards. Pay the immediate decimal contribution when a digit is placed:

$$r((\mathbf{m}, d), a = p) = d 10^p.$$

The nice thing about this rewards structure is that there's no need for additional calculation at the end of the game.

Objective. Maximize expected total reward over the five decisions:

$$\max_{\pi} \mathbb{E}_{\pi} \left[\sum_{t=0}^4 r(s_t, a_t) \right].$$

Problem 2

Finite-horizon DP formulation and solution

We have a finite horizon undiscounted MDP with 3 rounds. The state encodes the customer's excitement for the two products $s = (x_1, x_2) \in \{E, U\}^2$ at the start of each round and the action space $a \in \{1, 2\}$ represents the seller's recommendation.

One could consider adding the round number t to the state, but it is not necessary since the horizon is fixed and known.

Reward.

$$r(s, a) = \begin{cases} 1, & a = 1 \text{ and } x_1 = E, \\ 2, & a = 2 \text{ and } x_2 = E, \\ 0, & \text{otherwise.} \end{cases}$$

Transition model. At the end of the round, only the recommended product's excitement may flip:

$$\Pr(x'_1 = \neg x_1 \mid a = 1) = 0.1, \quad \Pr(x'_2 = \neg x_2 \mid a = 2) = 0.5,$$

and the other component stays unchanged.

Bellman Equation. Let $V_k(s)$ be the optimal expected revenue with k rounds remaining (so $V_0 \equiv 0$). Then for $k \geq 1$,

$$V_k(s) = \max_{a \in \{1, 2\}} [r(s, a) + \mathbf{E}[V_{k-1}(s') \mid s, a]].$$

Intuition / Strategy Product #2 offers a higher reward per purchase and over the long run, it will be available with the same frequency as product #1 (both products have symmetric flip probabilities). In an infinite horizon setting, the optimal policy is to always recommend product #2. However, in a finite horizon setting, it can be higher EV to recommend product #1 when product #2 is unavailable (state EU), and there are few turns left.

Backward induction. The state space is $\{EE, EU, UE, UU\}$. Let's calculate V_k and an optimal action $\pi_k^*(s)$.

$k = 1$ (one round left): At this point, we're just trying to maximize immediate reward. So, if the second product is available, we recommend it. If not, recommend the first product. If neither is available, the policy is indifferent.

s	EE	EU	UE	UU
$V_1(s)$	2	1	2	0
$\pi_1^*(s)$	2	1	2	\emptyset

$k = 2$: We now have to consider the final round as well. Our strategy will be similar to $k = 1$, but we have a preference to recommend product #2 because it has a higher reward.

s	EE	EU	UE	UU
$V_2(s)$	3.5	1.9	3.0	1.0
$\pi_2^*(s)$	2	1	2	2

$k = 3$: Our strategy is the same as $k = 2$

s	EE	EU	UE	UU
$V_3(s)$	4.7	2.81	4.0	2.0
$\pi_3^*(s)$	2	1	2	2

Optimal 3 round policy. From any state with $k \in \{2, 3\}$ rounds left, the same mapping applies:

$$EE \rightarrow 2, \quad EU \rightarrow 1, \quad UE \rightarrow 2, \quad UU \rightarrow 2,$$

with $k = 1$ agreeing except that UU is a tie. Starting from the given initial state EE (both eager), the optimal expected total revenue is

$$W_3(EE) = \boxed{4.7}.$$

Problem 3

Part (a)

Lets first restate the relationship between average reward and discounted value from our lecture slides:

$$\rho^\pi(s) = \lim_{\gamma \rightarrow 1} (1 - \gamma) V_\gamma^\pi(s)$$

Bias is defined as

$$h^\pi(s) = \lim_{T \rightarrow \infty} \mathbf{E} \left[\sum_{t=1}^T (r_t - \rho^\pi(s_t)) \mid s_1 = s \right],$$

and the slides show

$$h^\pi(s) - h^\pi(s') = \lim_{\gamma \rightarrow 1} (V_\gamma^\pi(s) - V_\gamma^\pi(s')).$$

There are two states s_1, s_2 ; in each, two actions a_1, a_2 with deterministic transitions and rewards:

$$\begin{array}{ll} s_1 \xrightarrow[a_1]{\$4} s_1, & s_1 \xrightarrow[a_2]{\$3} s_2, \\ s_2 \xrightarrow[a_1]{\$7} s_2, & s_2 \xrightarrow[a_2]{\$5} s_1. \end{array}$$

I will denote the four deterministic stationary policies by π^{ij} where i is the action in s_1 and j is the action in s_2 . So for example, π^{12} will loop in s_1 and transition from s_2 to s_1 .

Lets start in state s_1 and calculate the discounted values for each policy.

$$V_\gamma^{\pi^{11}}(s_1) = 4 + 4\gamma + 4\gamma^2 + \dots = \frac{4}{1-\gamma}$$

$$V_\gamma^{\pi^{12}}(s_1) = 4 + 4\gamma + 4\gamma^2 + \dots = \frac{4}{1-\gamma}$$

$$V_\gamma^{\pi^{21}}(s_1) = 3 + 7\gamma + 7\gamma^2 + \dots = 3 + \gamma \frac{7}{1-\gamma}$$

$$V_\gamma^{\pi^{22}}(s_1) = 3 + 5\gamma + 3\gamma^2 + 5\gamma^3 + \dots = \frac{3+5\gamma}{1-\gamma^2}$$

Lets move to state s_2

$$V_\gamma^{\pi^{11}}(s_2) = 7 + 7\gamma + 7\gamma^2 + \dots = \frac{7}{1-\gamma}$$

$$V_\gamma^{\pi^{12}}(s_2) = 5 + 4\gamma + 4\gamma^2 + \dots = 5 + \gamma \frac{4}{1-\gamma}$$

$$V_\gamma^{\pi^{21}}(s_2) = 7 + 7\gamma + 7\gamma^2 + \dots = \frac{7}{1-\gamma}$$

$$V_\gamma^{\pi^{22}}(s_2) = 5 + 3\gamma + 5\gamma^2 + 3\gamma^3 + \dots = \frac{5+3\gamma}{1-\gamma^2}$$

Moving on to average rewards, we can see that π^{21} is the only policy that yields an average reward of 7 in both states. π^{11} yields 4 in s_1 and 7 in s_2 . π^{12} and π^{22} both yield 4 in both states.

Part (b)

(1) -1 discount optimal (average-reward optimal): At s_1 , the only policy that maximizes average reward is π^{21} . At s_2 , both π^{11} and π^{21} maximize average reward. However, to be optimal at both states simultaneously, we must choose π^{21} . So the only -1 discount optimal policy is π^{21} .

(2) 0 discount optimal (bias optimal): Apologies in advance for the brevity, but we can quickly see that at s_2 , both π^{11} and π^{21} maximize the gain, regardless of discount applied. We'll conclude that π^{21} is the optimal 0-discount policy for the same reason as above.

The situation is more interesting at s_1 . Depending on the discount factor, it might be better to choose π^{11} or π^{21} . Lets setup the inequality:

$$\begin{aligned}
 V_{\gamma}^{\pi^{21}}(s_1) - V_{\gamma}^{\pi^{11}}(s_1) &\geq 0 \\
 3 + \gamma \frac{7}{1-\gamma} - \frac{4}{1-\gamma} &\geq 0 \\
 \frac{3(1-\gamma) + 7\gamma - 4}{1-\gamma} &\geq 0 \\
 3 - 3\gamma + 7\gamma - 4 &\geq 0 \\
 4\gamma - 1 &\geq 0 \\
 \gamma &\geq 0.25
 \end{aligned}$$

So for $\gamma \geq 0.25$, π^{21} is optimal at s_1 . For $\gamma < 0.25$, π^{11} is optimal at s_1 .

(3) 1 and ∞ discount optimal: The same logic applies as in (2). π^{21} is optimal at s_2 for all γ . At s_1 , π^{21} is optimal for $\gamma \geq 0.25$. So the only 1-discount optimal policy is π^{21} .

Part (c)

We have shown that when $\gamma \geq 0.25$, π^{21} is n-discount optimal for all $n \geq -1$. When $\gamma < 0.25$, π^{11} is optimal at s_1 but not at s_2 . So, our blackwell optimal policy is π^{21} , and the smallest discount factor that suffices is $\gamma^*(s_1) = 0.25$.

Blackwell optimal policy = π^{21} , $\gamma^*(s_1) = 0.25$, $\gamma^*(s_2) = 0$.
--